

OpenFlow and beyond on the Cisco Catalyst 9000

Atri Indiresan
atri@cisco.com
Cisco Systems
FAUCETCon 2019
October 23, 2019



Agenda

- **Catalyst 9K OpenFlow**
- **OpenFlow pipeline abstraction**
- **TFM and flexible pipelines**
- **Beyond OpenFlow: hybrid behavior**
 - **LLDP and PoE**
 - **MacSec and MKA**

Catalyst 9K OpenFlow

Cisco Catalyst 9000 enables Faucet OpenFlow

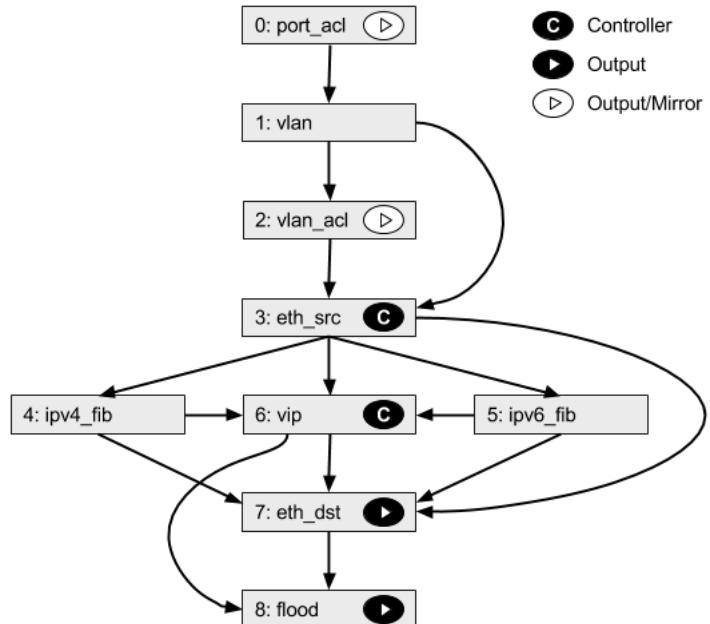


Built for
Programmability



Programmability

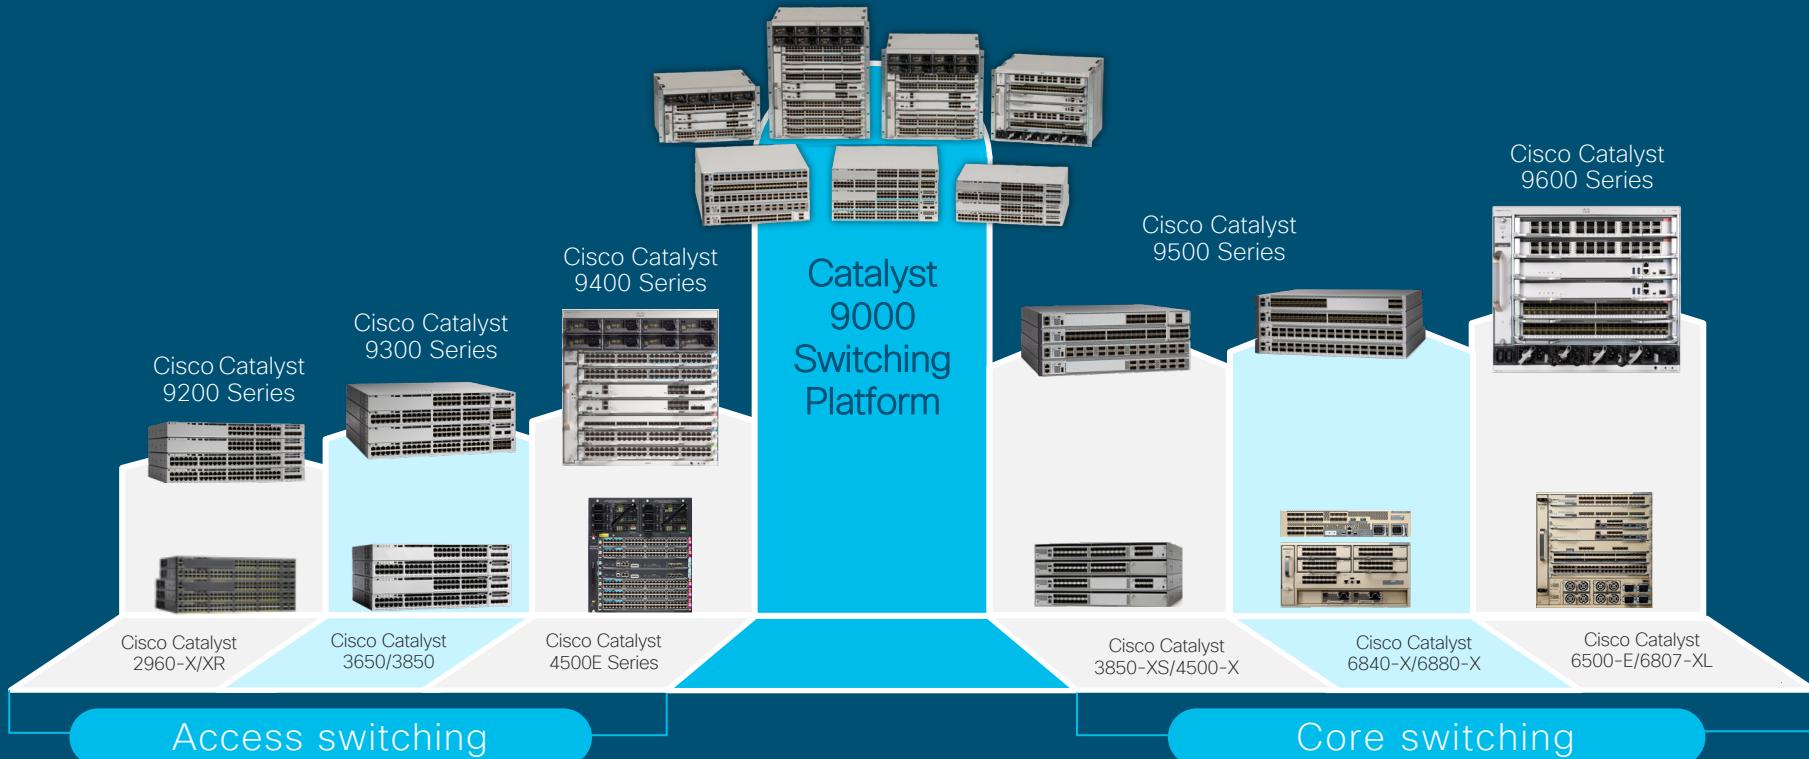
OpenConfig, GNMI, GNOI



Catalyst 9K OpenFlow in the wild

- General availability since IOS-XE 16.9.1 (August 2018)
 - Major release 3 times a year
- 4000+ fixed and modular switches deployed in enterprise at 40+ locations worldwide
 - Faucet and proprietary controllers
 - Hybrid features like LLDP for PoE utilized in production
 - Active feature development in coordination with customers

Cisco Switching Portfolio



Catalyst 9K OpenFlow Summary

- Powered by UADP
 - Programmable pipeline
 - Flexible match and actions
- Faucet and TFM compatible
- Software enhancements for features and scale

Platform	Type	Number of Tables	Number of Priority Levels	Number of Flows
9300	Fixed access, copper	16	128	9K
9400	Modular access, copper	16	128	27K
9500	Fixed agg/core, optical	16	128	27K
9500H	Fixed agg/core, optical, high performance	16	128	27K



Cisco Catalyst 9300 Series

New generation of fixed access

2.5G at the price of 1G

40G at the price of 10G

Only stackable switch with
8x 10G and 2x 25G uplinks



Highest 2.5G, 5G, Multigigabit density in the industry

Multigigabit + 2.5G

12 ports Multigigabit + 36 ports 2.5G

1G data

24 ports 1G

1G Cisco UPOE/PoE+

48 ports 1G

Multigigabit Cisco UPOE

48 ports Multigigabit

Cisco Catalyst 9000 leadership

UADP 2.0

Cisco IOS® XE Software

SD-Access

x86 CPU and containers

Encrypted Traffic Analytics (ETA)*

AES-256/MACsec-256

Trustworthy systems

Cisco StackWise® Virtual

IEEE1588 and AVB

NBAR2

Perpetual/Fast PoE

Model-driven programmability

Patching/GIR

Streaming telemetry

Modular fans

Modular uplinks

Higher-efficiency AC and DC power supplies



8x 10G

2x 40G

NEW

4x Multigigabit

NEW

4x 1G

2x 25G

NEW

315W AC

NEW

715W AC/DC*

NEW

1100W AC

Cisco Catalyst 9400 Series

New generation of modular access



Supervisor
• Sup-1: 80G per slot access optimized
• Sup-1XL*: 120G per slot core optimized

Access line cards
• 24x Multigigabit + 24x Cisco UPOE®
• 48x Cisco UPOE
• 48x PoE+, 48x data

Core line cards
• 24x 10G SFP+
• 48x 1G SFP
• 24x 1G SFP

Power supply
• 3200W AC
• 3200W DC*
• 2100W AC

Cisco® Catalyst® 9000 leadership

UADP 2.0
Open Cisco IOS® XE
SD-Access
x86 CPU and containers
Encrypted Traffic Analytics*
256-bit MACsec*
Trustworthy systems
Cisco StackWise® Virtual*
IEEE1588 and AVB*
NBAR2
Model-driven programmability
Patching and GIR
Streaming telemetry

Cisco Catalyst 9500 Series

New generation of purpose-built fixed core/aggregation

UADP 2.0



UADP 3.0



- 40G at the price of 10G
- Industry's first 40G enterprise switch
- Optimized 10G switch for midsize backbone



- 1TB SSD storage
- Industry's first 100G enterprise switch
- 8x buffering vs. the competition

Extending Cisco® Catalyst® 4000 and 6000 Series leadership in fixed core

13x throughput (3.2 Tbps)

6x performance (1 Bpps)

No oversubscription

8x 40G density

Pluggable SSD storage

USB 3.0

4x memory and flash

2x CPU cores

Customizable templates

Cisco StackWise® Virtual



Modular fans



Modular uplinks



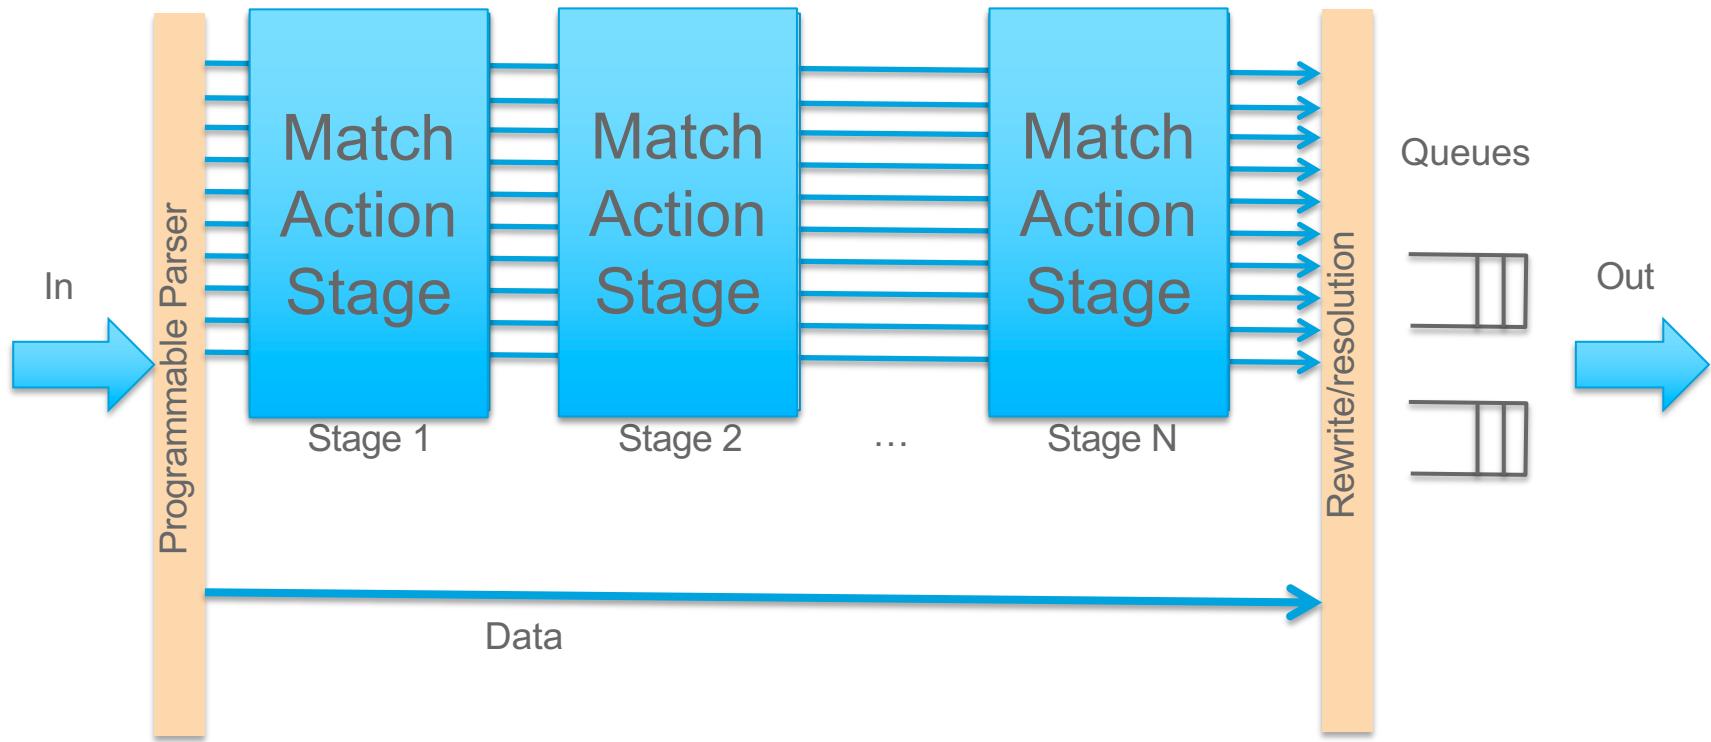
Modular power supplies



Storage for application hosting

OpenFlow Table Feature Message and flexible pipelines

Match-Action Table Pipeline

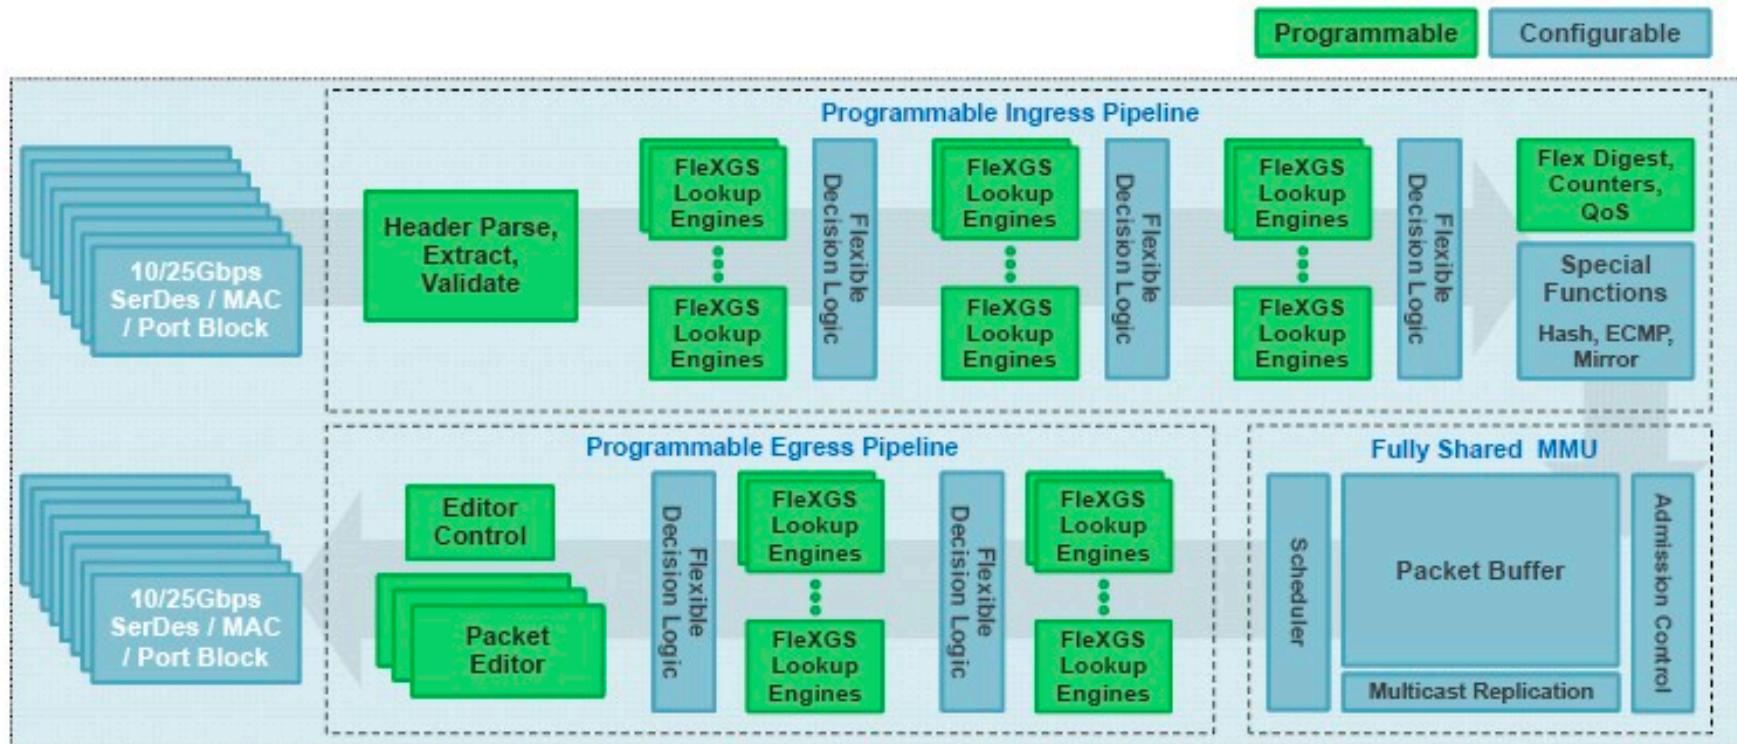


What's in your pipeline?

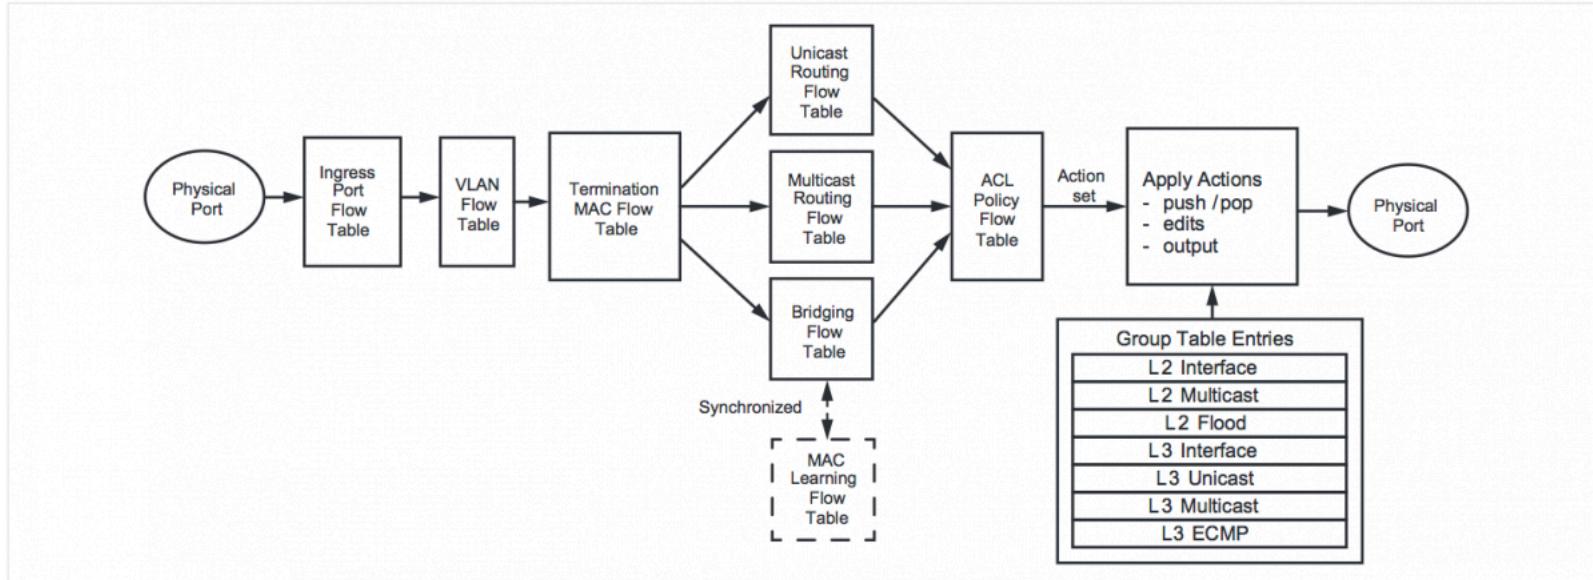
- Fixed vs Flexible
- Number of pipeline stages
- Operations associated with pipeline stages

With a sufficient number of recirculations, a fixed pipeline could emulate a flexible pipeline

Broadcom Trident 3 pipeline



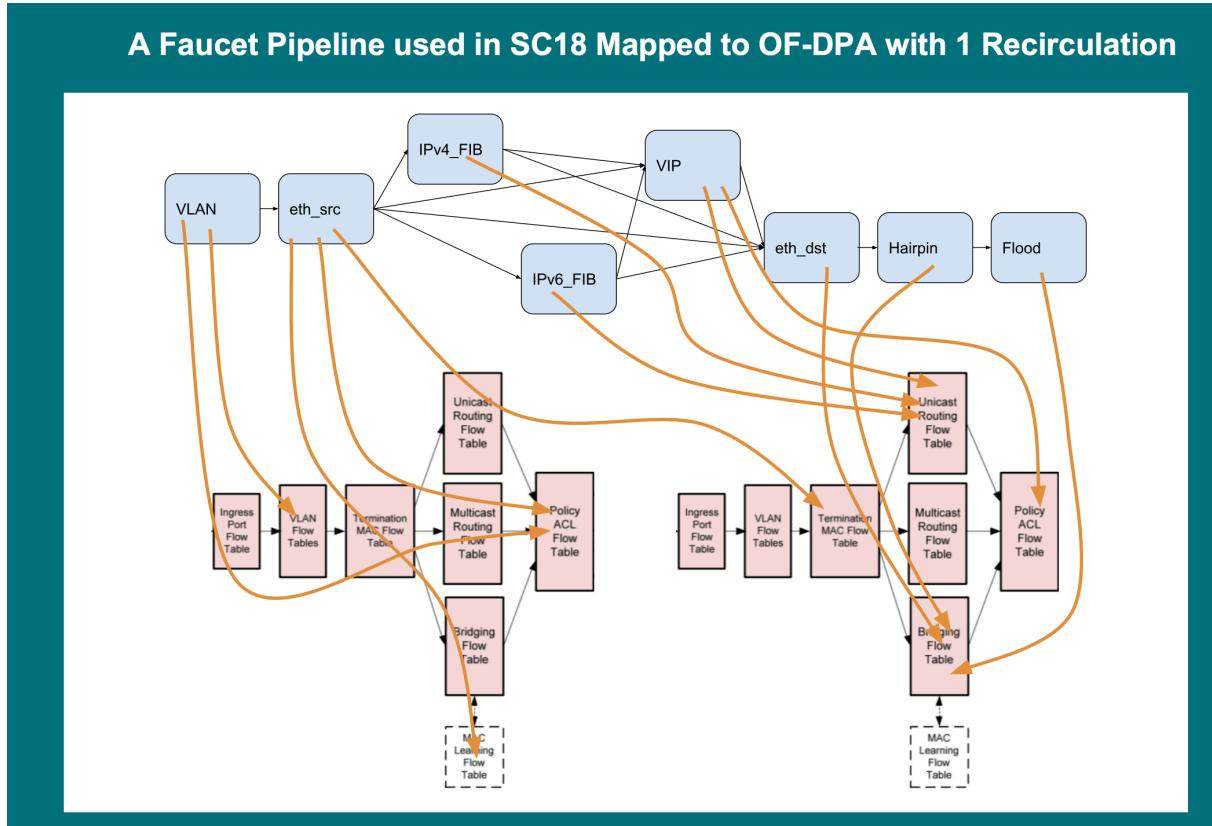
Broadcom abstract switch for OF-DPA



Abstract Switch Pipeline for Bridging and Routing

Example of Faucet with OF-DPA

Chris Lorier: <https://wand.net.nz/~cml16/poster.pdf>



OpenFlow Table Feature Message

- Basic parameters
 - Number of tables
 - Table size
- Table parameters
 - Match fields
 - Exact, wildcard, maskable
 - Actions
 - Apply actions
 - Forwarding – output, drop, punt
 - GoTo table

Outline of TFM algorithms

- Assign OpenFlow table to pipeline stage
 - Get next pipeline stage that matches the table's capabilities
 - Size, type of memory, matches and actions
 - If last pipeline stage, next pipeline stage is after recirculation
 - Some apply actions may force a recirculation
- Flow programming (add, delete, update)
 - Get pipeline stage associated with table
 - Program matches and actions

Beyond OpenFlow: Hybrid behavior

LLDP and PoE
MacSec and MKA

Native vs Hybrid behavior

Native behavior is fully specified by the OpenFlow pipeline and data plane

Hybrid behavior requires switch capabilities beyond what is specified by OpenFlow

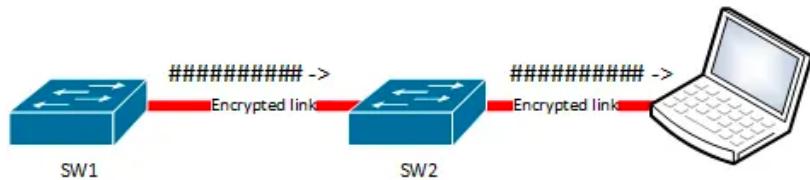
OpenFlow can redirect packets to the switch CPU with “action=output:LOCAL”

LLDP MED and Power over Ethernet

Media Endpoint Discovery extension of LLDP

- Auto-discovery of LAN policies
- Device location discovery
- Inventory management
- Extended and automated power management of *Power over Ethernet (PoE)* end points
 - How does the controller manage power levels of end-points?
 - Is it better for the switch to do this?

MacSec Key Agreement (MKA) – 802.1AE



- Key negotiation
- Rekeying
- EAPOL BPDUs



What's next?

Future directions – devices and controllers

New features and capabilities must be driven by deployment needs and Faucet requirements

- Network architecture and solutions
 - Multipath - L2 and L3
 - Segmentation
 - Identity
 - Policy

Catalyst 9K OpenFlow team

Atri Indiresan

Christophe Rene

Imran Mohamed

Manas Pati

Mel Tsai

Mohanraj K

Nilesh Ashok Inamdar

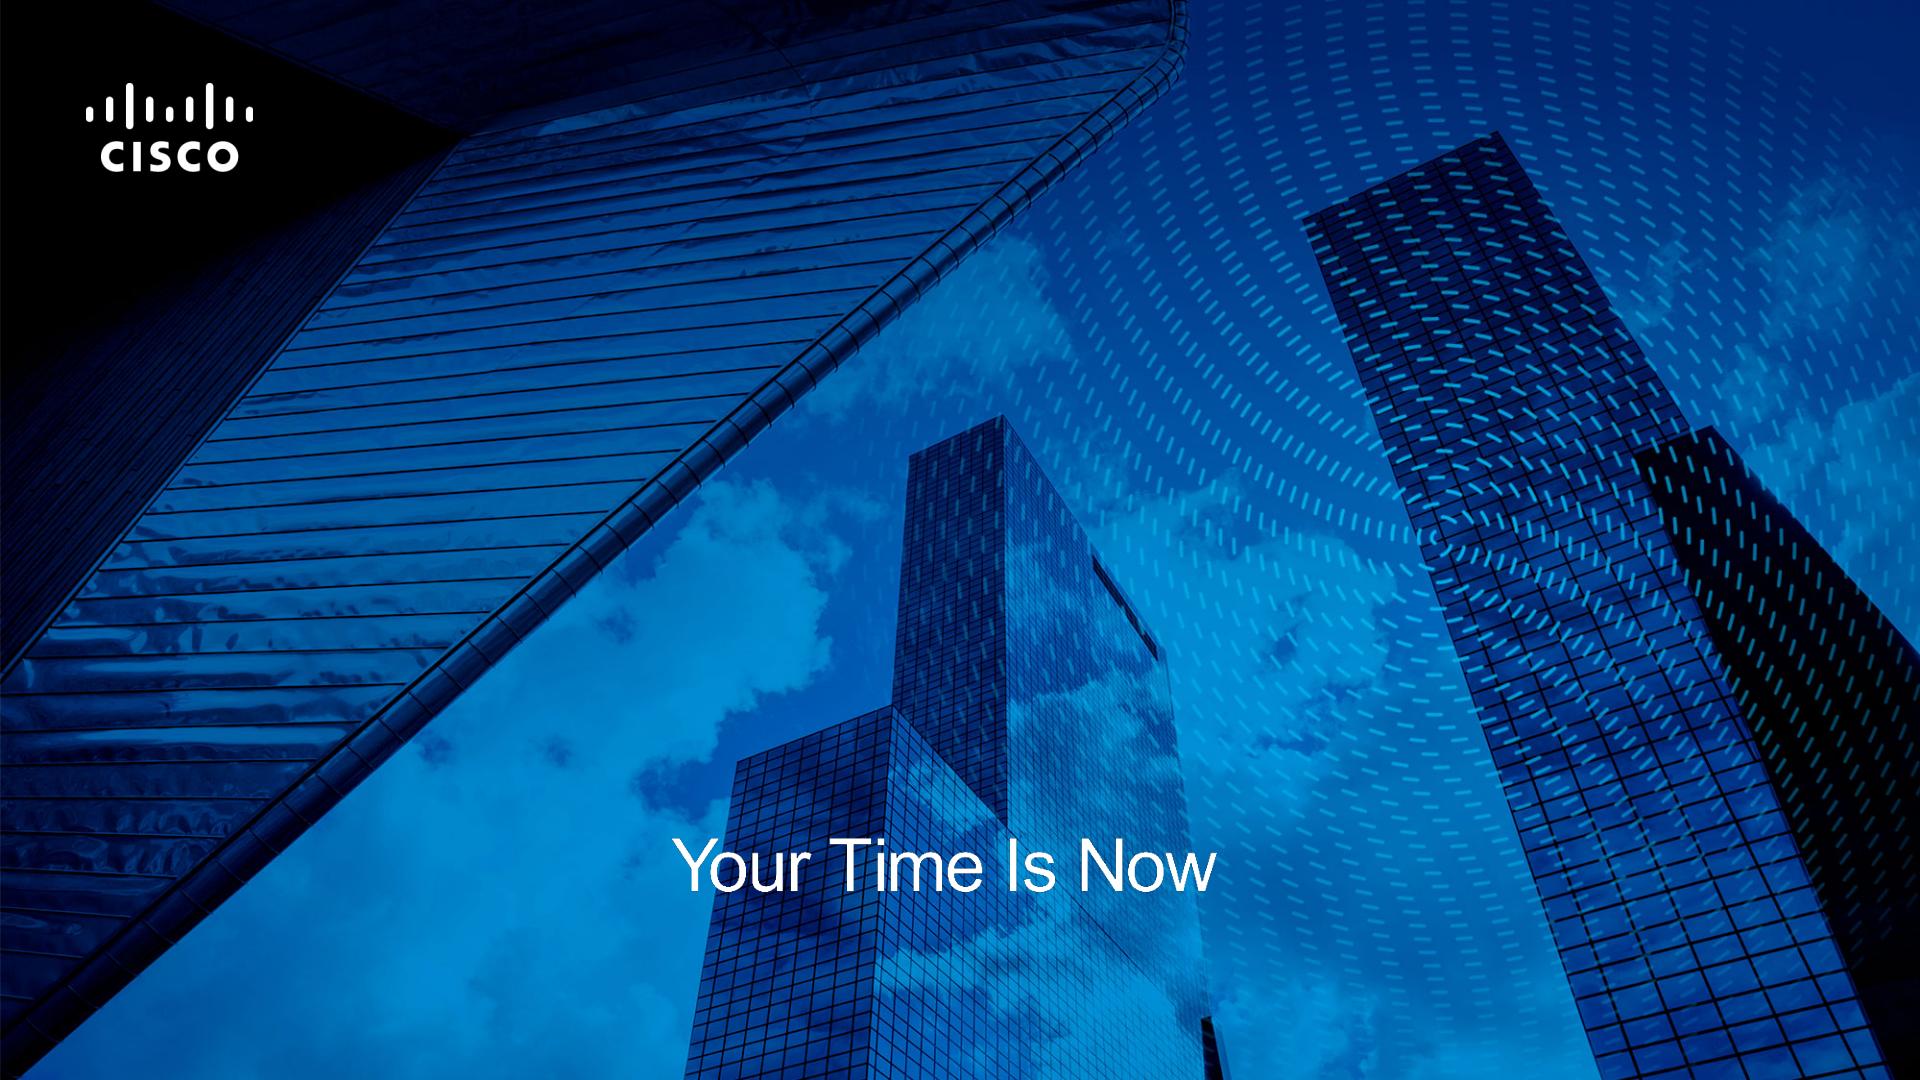
Nirupa Gnanasekaran

Orkhan Hasanli

Stephen Wang

Suneetha Maddineni

Q & A

A dark blue-toned abstract background featuring several modern skyscrapers with glass and steel facades. The buildings are set against a backdrop of light blue and white clouds, creating a sense of depth and motion. The overall composition is dynamic, with the buildings appearing to move towards the viewer.

Your Time Is Now